Paris, samedi 14 décembre 2019

Ecole Supérieure de Génie Informatique

Merci à l'ESGI, aux communautés et aux sponsors

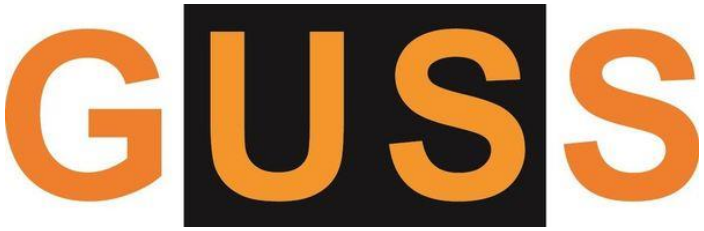

Local sponsors

Global AI Bootcamp

Samedi 14 décembre 2019 – Paris

Custom Vision : projet de reconnaissance de l'équipement de protection

Le projet de contrôle de tenue, qui permet d'améliorer les conditions de travail dans les usines du client. Le projet comprend 2 parties : détection de visage (réalisée avec OpenCV) et classification de tenue (réalisée avec Computer Vision API de Microsoft)

Alibek JAKUPOV (EXPERTIME)

Microsoft MVP (AI)

Global AI Community

15h-16h / Track vert

Use cases

Explanation

Protective Clothing detection

Requirements: Win 10 application connected to USB camera and detecting multiple operators simultaneously

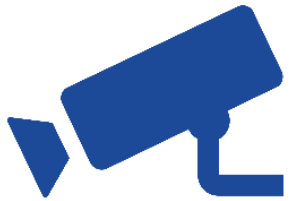

Possible solutions

Deep Learning

- CNTK; TensorFlow; Keras etc...

Cascade Classifier

- OpenCV
- Bibliothèque open source spécialisée dans le traitement d'image en temps réel.

Transfer Learning

- Custom Vision API
- API permettant de créer, déployer et améliorer des classificateurs d'image personnalisés.
- Utilise un algorithme d'apprentissage automatique pour classifier les images.

Deep Learning

Use N-series VMs for GPU acceleration

Deep Learning

Use N-series VMs for GPU acceleration

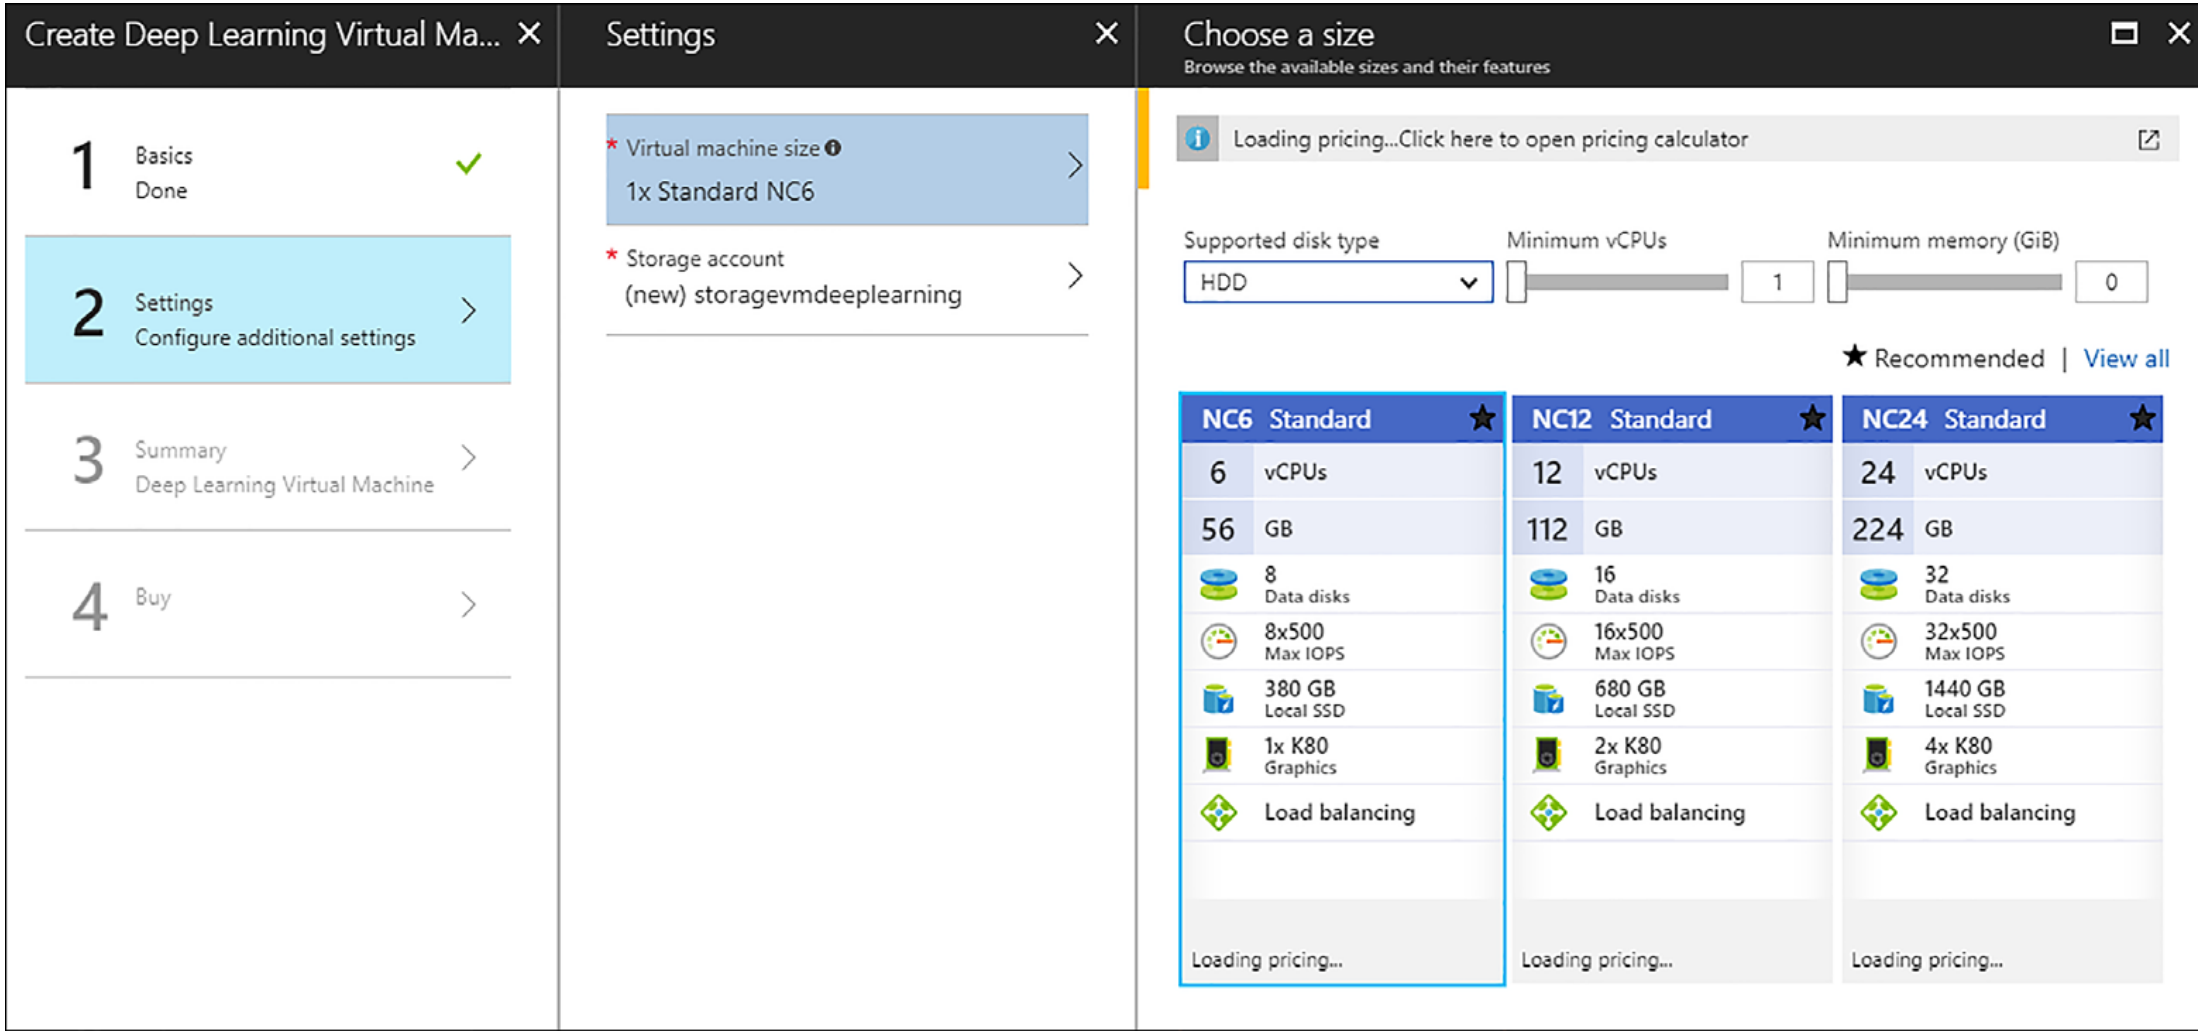

Deep Learning

Use N-series VMs for GPU acceleration

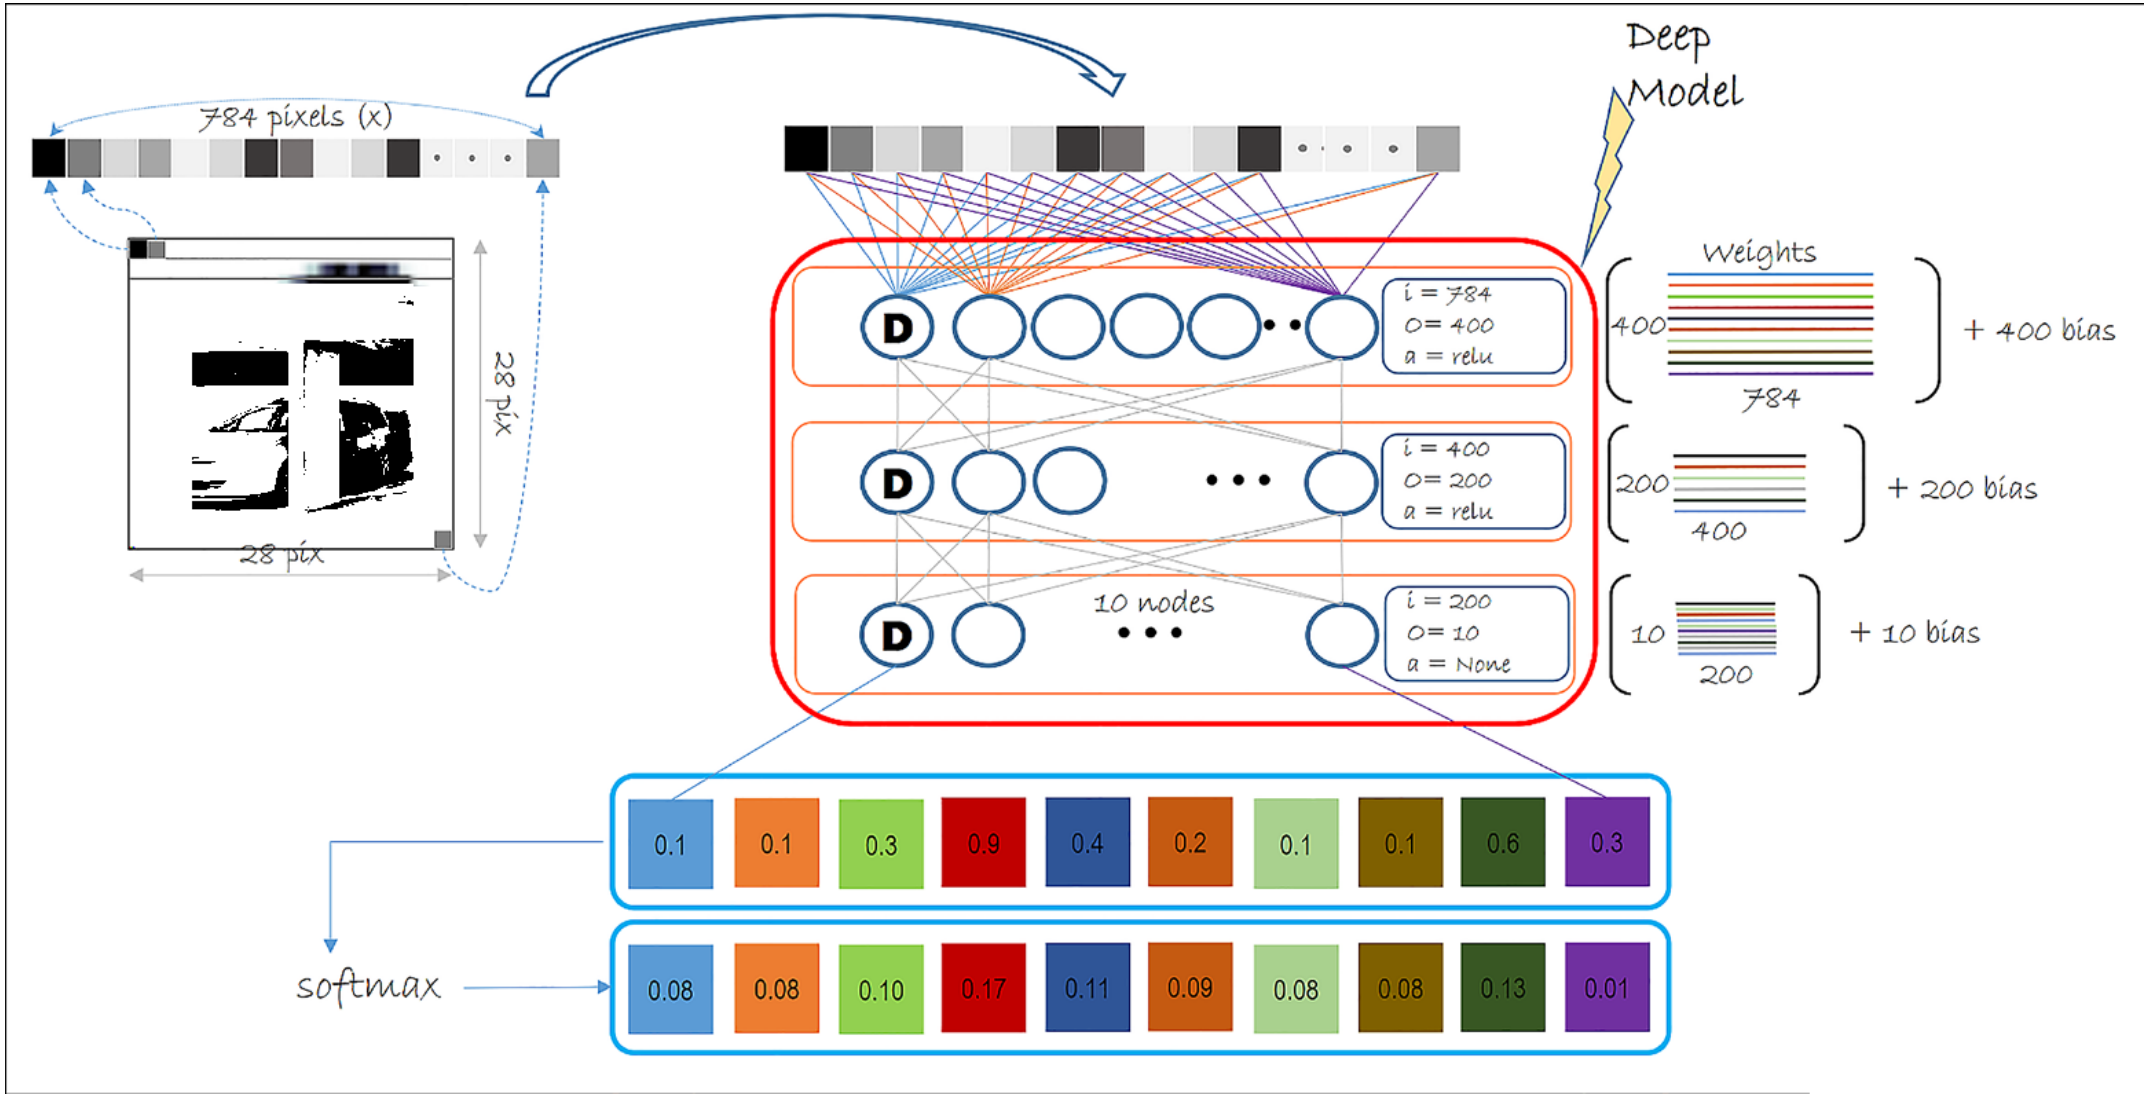

Deep Learning

Use N-series VMs for GPU acceleration

```
0 - Train Loss: 2.403, Test Accuracy: 0.197
500 - Train Loss: 0.230, Test Accuracy: 0.949
1000 - Train Loss: 0.058, Test Accuracy: 0.966
1500 - Train Loss: 0.037, Test Accuracy: 0.968
2000 - Train Loss: 0.072, Test Accuracy: 0.972
2500 - Train Loss: 0.064, Test Accuracy: 0.972
3000 - Train Loss: 0.058, Test Accuracy: 0.976
3500 - Train Loss: 0.036, Test Accuracy: 0.977
4000 - Train Loss: 0.034, Test Accuracy: 0.975
4500 - Train Loss: 0.004, Test Accuracy: 0.976
5000 - Train Loss: 0.059, Test Accuracy: 0.975
5500 - Train Loss: 0.010, Test Accuracy: 0.980
6000 - Train Loss: 0.018, Test Accuracy: 0.981
```

Deep Learning

Pros:
Flexibility
Precision

Cons:
Time
Complexity
Large data set

Cascade Classifier

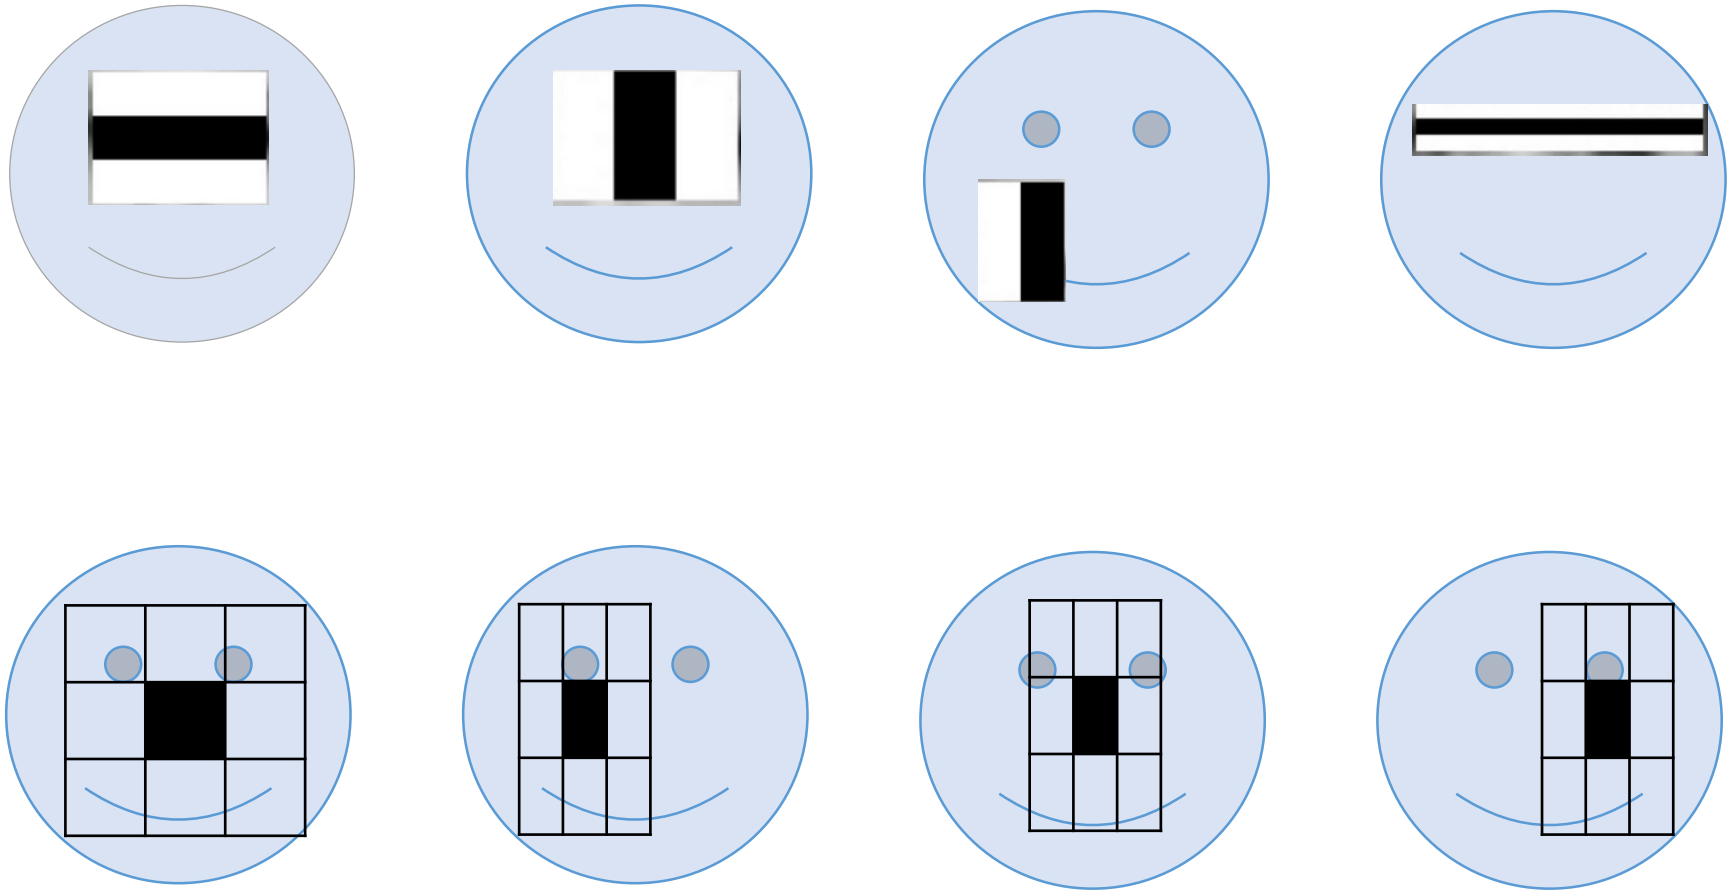

Cascade Classifier

Pros:
Speed
Open-source

Cons:
Precision
Programming skills required
Image processing skills required
Large training set needed
No user-friendly interface

Transfer Learning

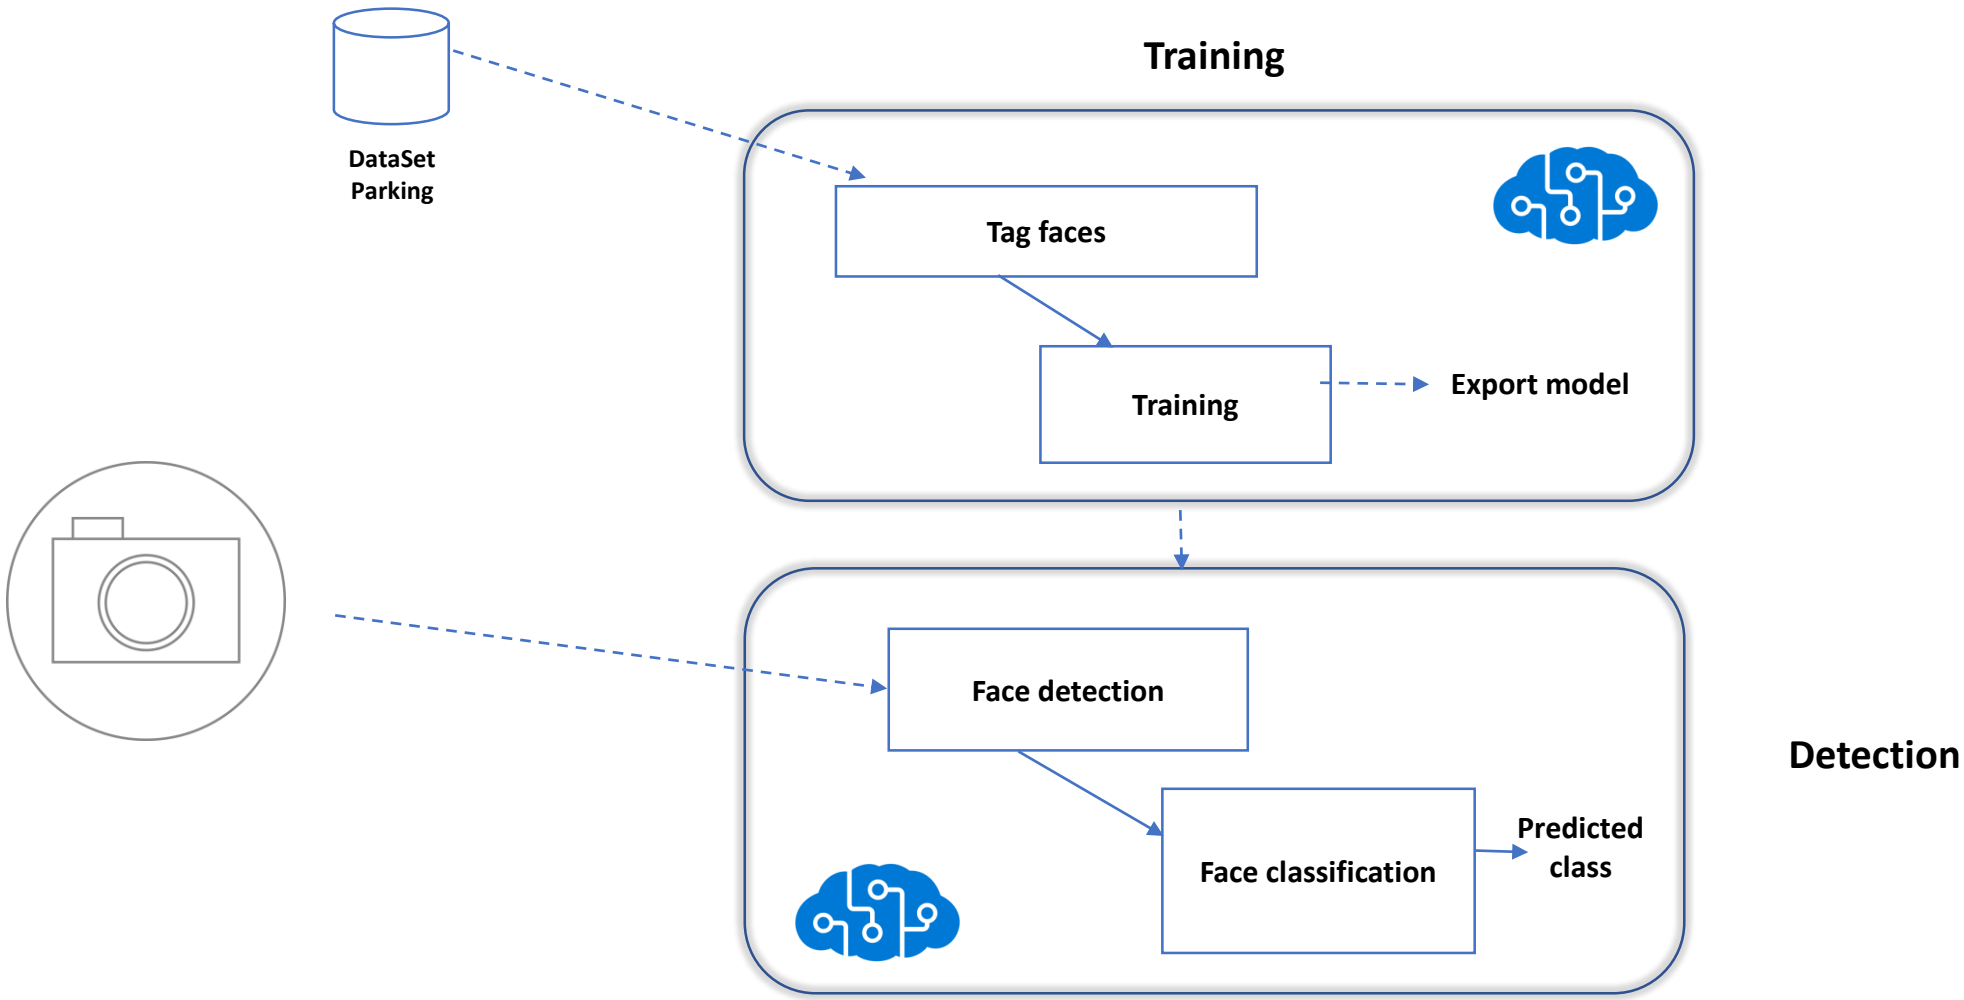

Transfer Learning

Pros:
Speed
High precision
Small dataset required
User-friendly interface

Cons:
Certain limits in model complexity

Transfer Learning

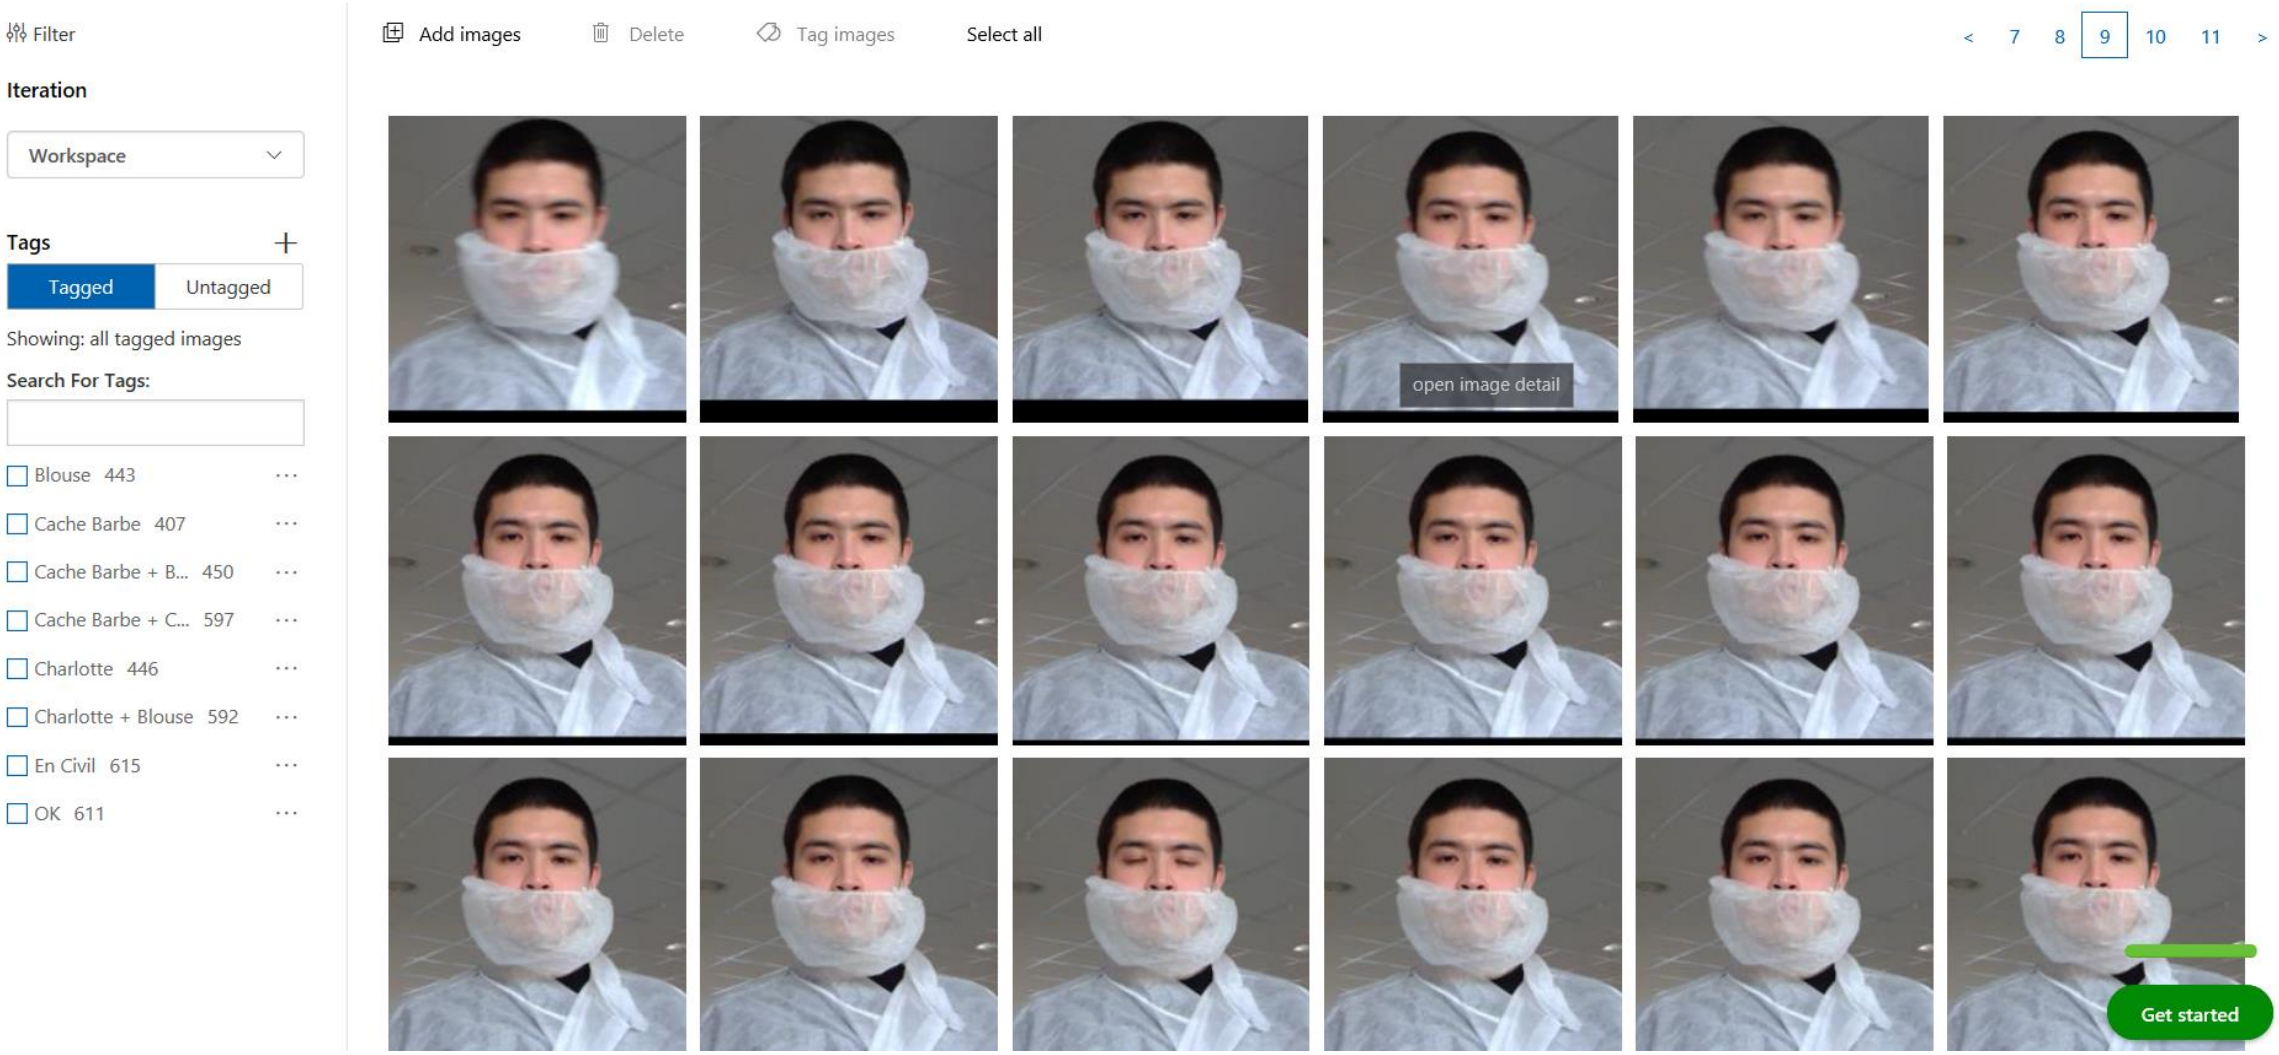

Trainset Generation

Dataset Generation

OK

Blouse

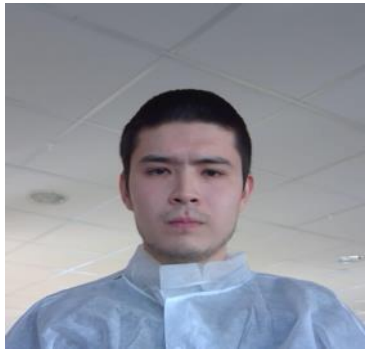

Cache Barbe + Blouse

En civil

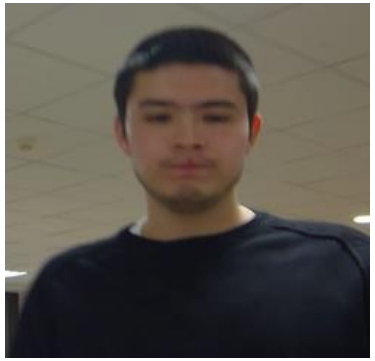

Charlotte

Blouse

Blouse + Charlotte

Cache Barbe

Detection of protective clothing

Hairnet Detection

Lab Coat Detection

Safety mask Detection

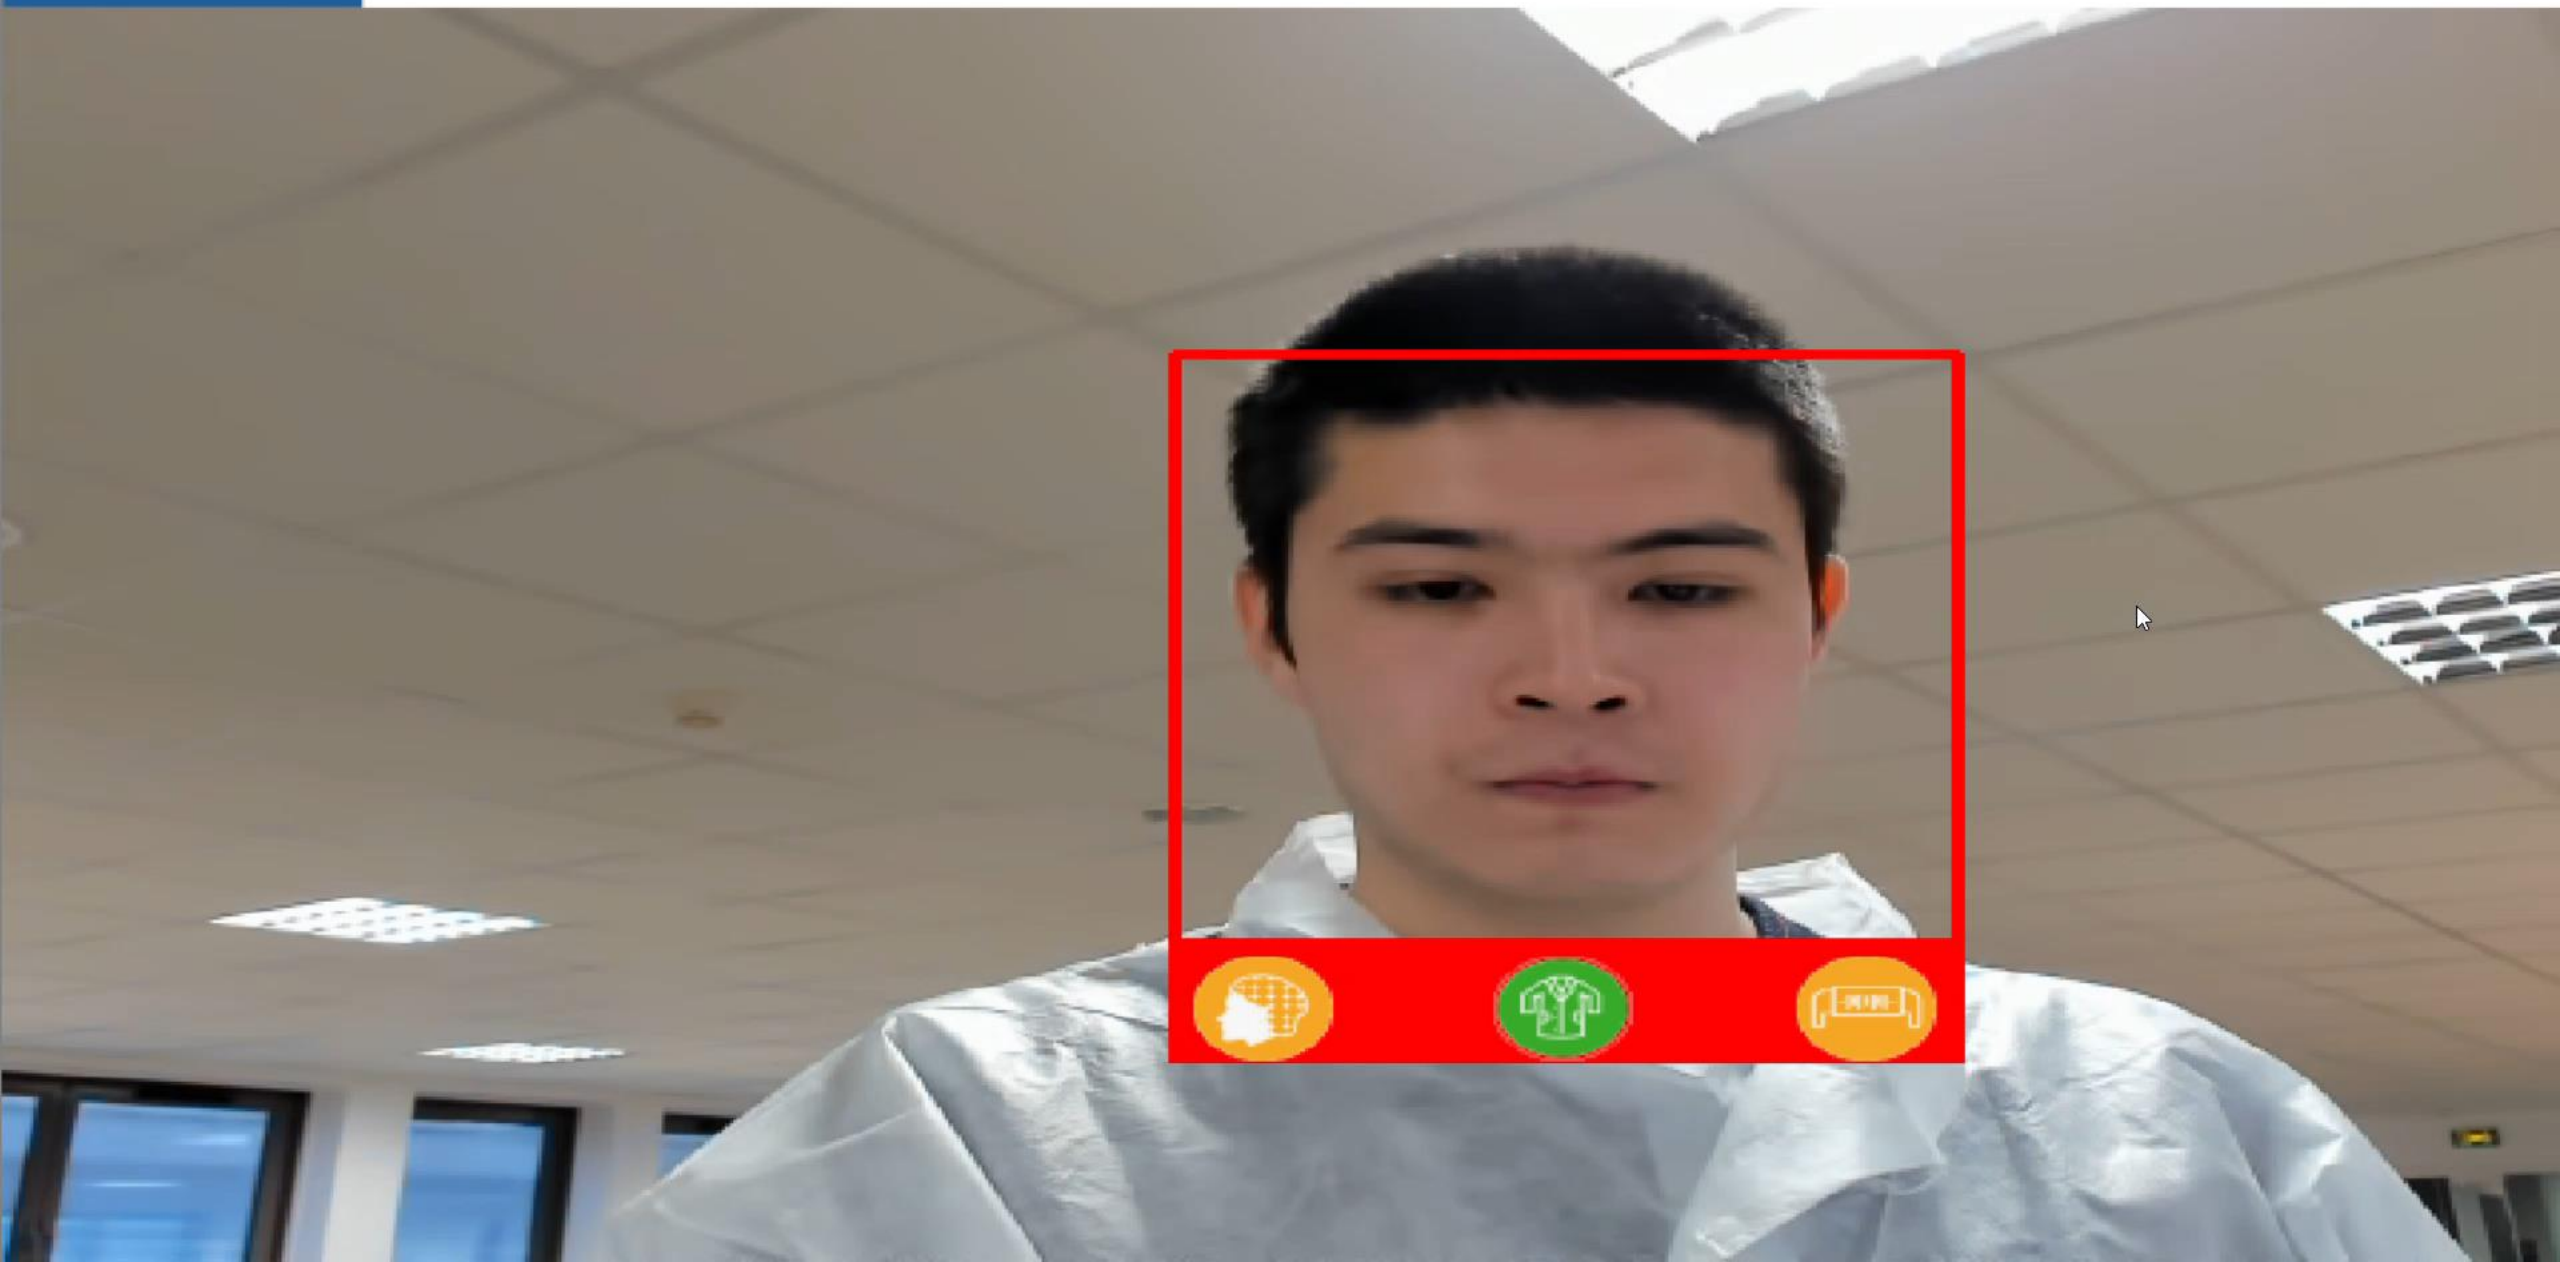

Important points

Problems

- True positives
- Distance/light condition
- Usage scenarios
- Computing speed

Tips and Tricks

Tips and Tricks

- Using Python 3.6 on Windows
- Cold-light
- Define optimal camera distance
- Screen size

Discussions

Alternative and complementary bricks and services

- Improve face detection
- Improve processing speed
- Improve UX
- Improve quality

Improve face detection

August 2017, OpenCV 3.3 was officially released, bringing it with it a highly improved “deep neural networks” (dnn) module.

Improve face detection

However, what most OpenCV users do not know is that Rybnikov has included a more accurate, deep learning-based face detector included in the official release of OpenCV

https://www.pyimagesearch.com/2018/02/26/face-detection-with-opencv-and-deep-learning/

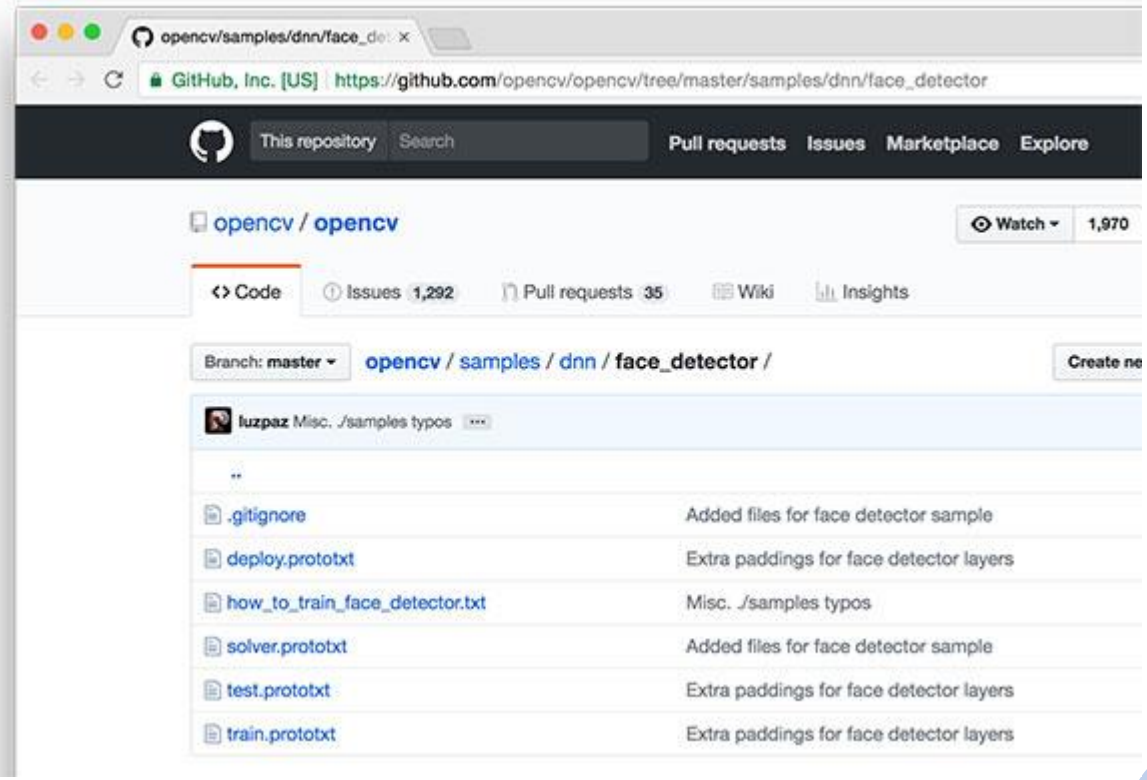

Improve face detection

- Caffe prototxt files
- Caffe model weight files

```
# load serialized model from disk
net = cv2.dnn.readNetFromCaffe(prototxt, model)

videoCapture = cv2.VideoCapture(0)

cv2.namedWindow("Face detect", cv2.WND_PROP_FULLSCREEN)
cv2.setWindowProperty(
    "Face detect", cv2.WND_PROP_FULLSCREEN, cv2.WINDOW_FULLSCREEN)

while (videoCapture.isOpened()):
    ret, frame = videoCapture.read()

    image = cv2.flip(frame, 1)

    # get image dimensions
    h, w = image.shape[:2]
    # resize image
    resized_image = cv2.resize(image, (300, 300))
    blob = cv2.dnn.blobFromImage(resized_image, 1.0,
                                 (300, 300), (104.0, 177.0, 123.0))

    net.setInput(blob)
    detections = net.forward()

    for i in range(0, detections.shape[2]):
        confidence = detections[0, 0, i, 2]

        if confidence > 0.5:
            box = detections[0, 0, i, 3:7] * np.array([w, h, w, h])
            uX, uY, dX, dY = box.astype("int")
            cv2.rectangle(img=image, pt1=(uX, uY), pt2=(
                dX, dY), color=WHITE, thickness=LINE_TYPE)

    cv2.imshow('Face detect', image)
```

Improve processing speed

- Tensorflow model on Linux run faster
- Full screen works better

Improve UX

- Use Electron JS to launch python scripts
- Create local server for prediction
- Add additional steps for verification

Improve quality

- Split model into two for better focus
- Add more data

Visualize Graph

pip install tensorboard

tensorboard --logdir=log

TensorBoard 1.11.0 at http://<computer-name>:6006 (Press CTRL+C to quit)

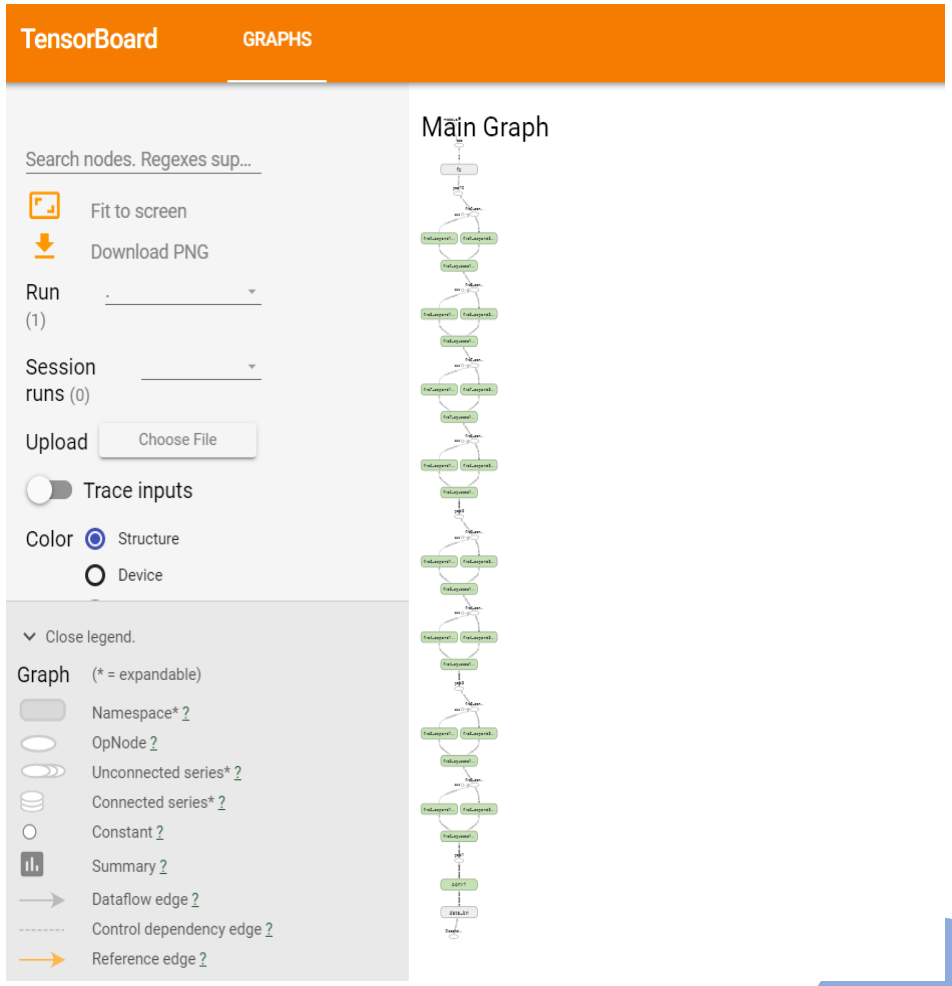

Edge Computing

GPU	**128-core Maxwell**
CPU	Quad-core ARM A57 @ 1.43 GHz
Memory	4 GB 64-bit LPDDR4 25.6 GB/s
Storage	microSD (not included)
Video Encode	4K @ 30 | 4x 1080p @ 30 | 9x 720p @ 30 (H.264/H.265)
Video Decode	4K @ 60 | 2x 4K @ 30 | 8x 1080p @ 30 | 18x 720p @ 30 (H.264/H.265)
Camera	1x MIPI CSI-2 DPHY lanes
Connectivity	Gigabit Ethernet, M.2 Key E
Display	HDMI 2.0 and eDP 1.4
USB	4x USB 3.0, USB 2.0 Micro-B
Others	GPIO, I2C, I2S, SPI, UART
Mechanical	69 mm x 45 mm, 260-pin edge connector

Edge Computing

Rewrite code on c++ using libjetson-inference library

Edge Computing

Edge Computing

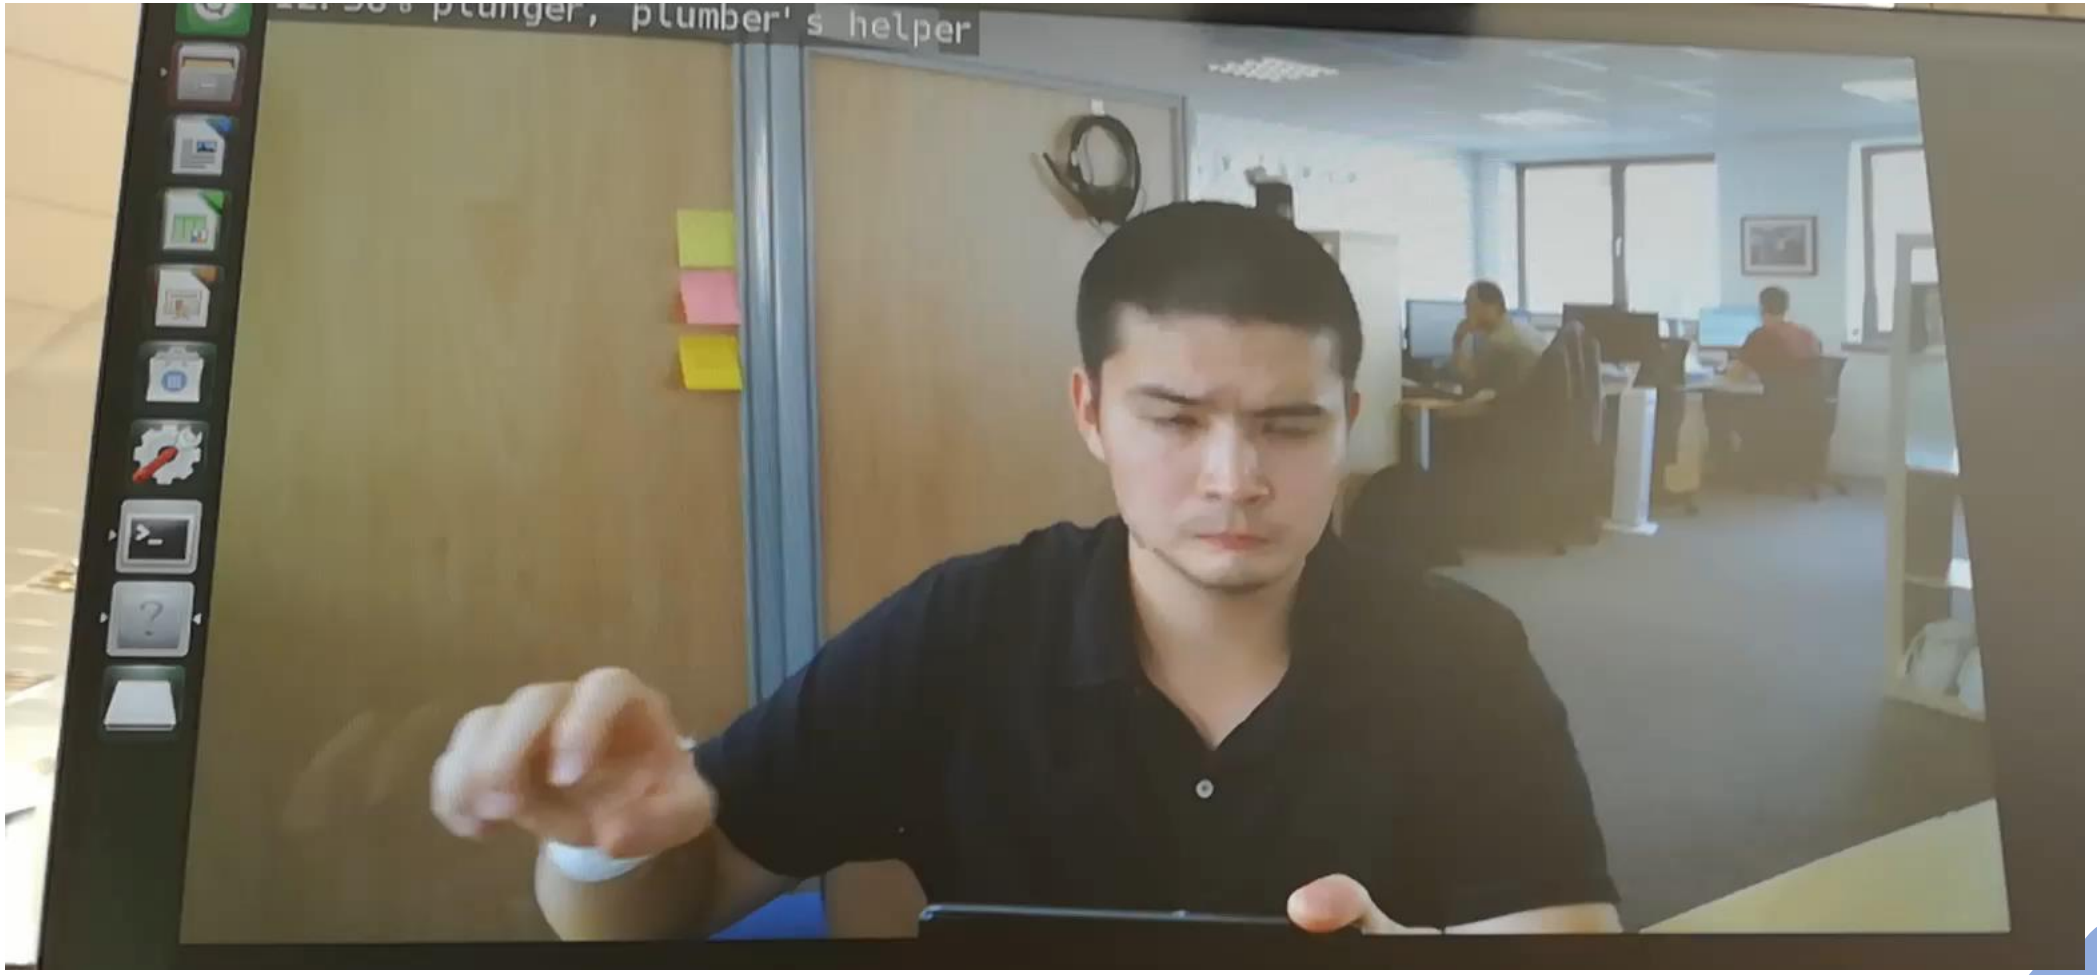

Edge Computing

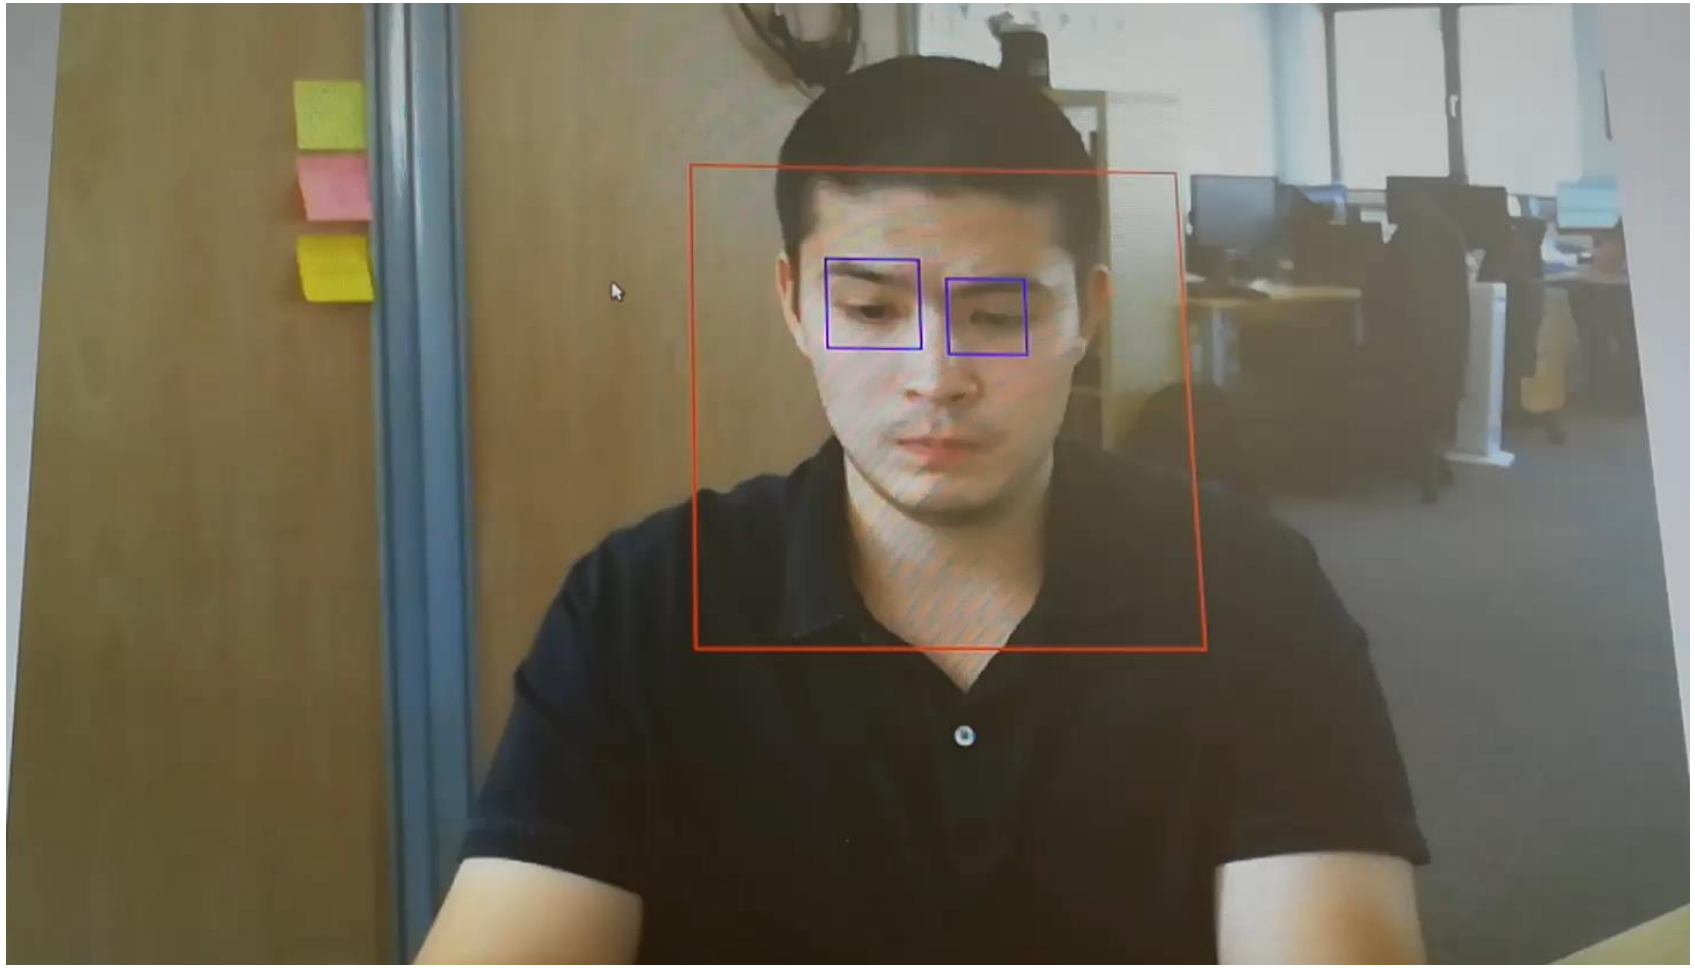

Paris, samedi 14 décembre 2019

Ecole Supérieure de Génie Informatique